


We taught sand to think, and we did it before we learned to say what we wanted it to do. There is a real, settled, and genuinely beautiful theory of what intelligence is: Bayes' theorem, Solomonoff's universal prior, and expected-utility maximization, synthesized as AIXI. And there are the messy, bounded systems we actually built, the large language models everyone now uses. Between the clean theory and the crude practice lies a gap, and that gap is where the safety of all of it will be decided. *On Intelligence and Its Specifications* builds both halves from the ground up, for the determined lay reader. It shows that the hardest problem is not making machines capable; it is saying what we want them to do. No hand-waving, and no doom. The mathematics earns every claim, and the conclusions are yours to draw.

On Intelligence and Its Specifications
Alex Towell

On Intelligence and Its Specifications



Alex Towell